



# JATIN

## DATA ENGINEER, 5+

### Summary:

Data Engineer with 5+ years of experience designing, building, and optimising production-grade cloud data warehouses, ELT pipelines, and data models across procurement, financial services, IP analytics, and healthcare domains. Deep practical experience with Snowflake (Streams, Tasks, Zero-Copy Cloning, Clustering Keys, Dynamic Data Masking, RBAC, Data Sharing), DBT (staging/intermediate/mart layers, incremental models, tests, macros), and Apache Airflow orchestration. Strong Python engineer with proven delivery across Azure and AWS cloud environments.

---

### SKILLS

**Snowflake:** Streams & Tasks, Zero-Copy Cloning, Clustering Keys, Materialized Views, Result Caching, Dynamic Data Masking, Row Access Policies, RBAC, Data Sharing, Snowpipe, Query

**DBT:** Optimization, Warehouse Cost Management, Auto-Suspend

**Data Modelling:** version-controlled SQL — dbt Core & dbt Cloud Star Schema, Snowflake Schema, Data Vault, SCD Type 2, fact & dimension tables, OBT, normalised & denormalised designs  
**ELT / ETL:** Python-based ingestion pipelines, REST API extraction,

### EXPERIENCE

#### WITQUALIS TECHNOLOGIES PVT LTD

(CURRENTLY)

As Data Engineer

#### OTHER ORGANIZATION

Different Roles

### PROJECTS

#### 1. Procurement Data Warehouse — Snowflake + DBT + Airflow (StatusNeo)

Stack: Snowflake · DBT · Apache Airflow · Python · SAP/Oracle ERP APIs · Git · Slack

▸ End-to-end procurement DWH: fact tables (POs, GRNs, Invoices) + dimension tables (vendors, cost centres, categories); full DBT model layer with staging → intermediate → mart, dbt tests on every model, Git version control.

▸ 100+ vendor scorecards: on-time delivery, rejection rates, payment terms compliance — live in Snowflake-backed reporting layer; 35% Snowflake cost reduction via auto-suspend + query optimisation.

▸ Airflow DAG orchestration: daily ERP ingestion + dbt runs + data quality checks; Slack alerting on failures; Snowflake RBAC ensuring role-based data visibility for procurement, finance & execs.

#### 2. Financial Data Migration — QuickBooks → Snowflake (ISSC)

Stack: Snowflake · Python · QuickBooks API · Apache Airflow · SCD Type 2 · SQL



schema mapping,  
currency conversions,  
fiscal year alignment,  
semi-structured data  
(JSON, XML, VARIANT,  
FLATTEN) Apache  
Airflow (DAGs, retry  
logic, SLA monitoring,  
Slack alerting), Azure  
Data Factory, dbt job  
scheduling

**Orchestration:** Apache  
Airflow (DAGs, retry  
logic, SLA monitoring,  
Slack alerting), Azure  
Data Factory, dbt job  
scheduling

**Data Quality:**  
Automated reconciliation  
frameworks, dbt tests,  
Great Expectations, audit  
trail frameworks, zero-  
tolerance validation  
scripts

**Cloud & Storage:**  
Azure (Databricks, ADLS  
Gen2, ADF, Key Vault) ·  
AWS (S3, Glue, EC2) ·  
Snowflake on  
Azure/AWS

**Python & SQL:** Python  
3, Pandas, PySpark,  
Snowflake SQL, stored  
procedures, views,  
window functions, CTEs,  
DuckDB, SQLAlchemy

**Governance &  
Security:** Dynamic Data  
Masking, Row Access  
Policies, RBAC role

- Zero data loss migration of GL, AP/AR, invoices, payroll & trial balance data into Snowflake financial DWH: SCD Type 2 for full audit trail; automated Python reconciliation framework achieving 100% balance accuracy.
- CFO-ready P&L, balance sheet & cash flow reporting layer: directly queryable in Snowflake; Airflow-orchestrated with full error handling, retry logic, and execution logging.

### 3. Global Patent Analytics Platform — Snowflake + DBT (Clarivate)

Stack: Snowflake · DBT · Apache Airflow · AWS S3 · Python · SQL

- Unified 50+ global PTO offices into single Snowflake DWH: Star Schema for inventors/assignees/IPC classifications; Snowflake Streams & Tasks for incremental ELT (72 hrs → 4 hrs freshness).
- 65% query performance improvement: via Clustering Keys, Materialized Views & Result Caching on billion-row datasets; ZeroCopy Cloning for dev/test; Row Access Policies + Dynamic Data Masking for tiered client licensing.
- DBT transformation layer: modular staging/intermediate/mart models, dbt tests, version-controlled SQL; Airflow orchestration with 99.8% uptime SLA.

### 4. Global Clinical Data Platform — Snowflake + Databricks + FHIR (Agilon Health)

Stack: Snowflake · Azure Databricks · PySpark · Apache NiFi · Python · AWS S3 · Azure ADLS · FHIR R4 · HL7

- Consolidated EHR/EMR data from Epic, Cerner & Meditech across US, Australia & Brazil into Snowflake clinical DWH: FHIR R4 pipelines, PHI masking, field-level encryption, HIPAA/LGPD compliance.
- PySpark Medallion pipelines (Bronze/Silver/Gold): millions of patient records; automated Python data quality framework saving 20+ hrs/week; patient risk scoring & population health dashboards on Snowflake datamart.



hierarchies, PHI masking,  
HIPAA compliance, audit  
logging, Unity Catalog

**Databases:** Snowflake,  
Azure Databricks (Delta  
Lake), PostgreSQL,  
MySQL, MongoDB,  
SQLite, SQL Server

**Version Control:** Git,  
GitHub, branching  
strategies, dbt project  
version control

## **EDUCATION**

B.Tech — Computer  
Science & Engineering |  
GL Bajaj Group of  
Institutions, Mathura |  
Aug 2018 – Jul 2022