

HIMANSHU

DATA ENGINEER, 15+

Summary: Specialist Data & AI Engineer with 15+ years of experience designing and deploying enterprise-scale data pipelines, machine learning workflows, and Generative AI platforms. Expert in Azure, Databricks, and GenAI (RAG, LLM), with a proven track record of transitioning complex legacy data ecosystems into modern Lakehouse architectures. Demonstrated success in realising \$180K–\$240K in annual cost savings.

SKILLS

AI & ML Engineering

RAG

LLM Orchestration

Agentic AI

NLP

Prompt Engineering

ML Ops

Data Engineering

PySpark

Spark SQL

Delta Lake

ETL/ELT

Informatica

Data Platforms

Azure Databricks

Azure Synapse

Azure Data Factory (ADF)

Azure Data Lakes

DataOps & DevOps

Lakehouse Design

CI/CD

Azure DevOps

PROJECTS

Project 1: Enterprise GenAI & Data Platform Transformation

Role: Senior Manager, Data & AI

Project Description:

Designed and implemented an enterprise-scale Generative AI and Data Platform modernization initiative focused on improving business access to data, modernizing legacy workloads, and enabling real-time analytics. The platform leveraged Azure, Databricks, Delta Lakehouse, and RAG-based GenAI capabilities to bridge the gap between structured and unstructured enterprise data while reducing operational costs and improving data processing performance.

Responsibilities:

- Architected an enterprise GenAI (RAG) platform improving data access efficiency by 35% and bridging the gap between vast unstructured data stores and actionable business insights.
- Led the migration of 20+ legacy data workloads and applications to Azure using ADF and Databricks, enhancing data deployment agility by 40%+.
- Built 30+ highly scalable data pipelines on Delta Lakehouse processing 2+ TB of daily enterprise data for near real-time analytics.
- Realised \$180K–\$240K in annual savings (20% reduction in cloud data spend) while improving workload processing performance by 30%.
- Managed and mentored a team of 6–8 data and AI engineers, guiding data strategy, ML model productionalization, and continuous delivery.

Project 2: Azure Data Lake & DataOps Modernization

Role: Lead Data Engineer

Project Description:

Built a cloud-native Azure Data Platform from the ground up to support

EDUCATION

Bachelor of Engineering (B.E.)

**Electronics & Telecommunication
Engineering**

R.T.M. Nagpur University
2010

CERTIFICATION

Databricks GenAI Engineer

Azure Data Engineer Associate (DP-
203)

AWS Certified Solutions Architect

enterprise analytics, reporting, and large-scale data processing requirements. The project focused on establishing reusable ingestion frameworks, improving data pipeline efficiency, and implementing DevOps practices for streamlined deployments and operational excellence.

Responsibilities:

- Built foundational Azure data lake and cloud infrastructure from scratch, achieving ~\$100K+ in cost savings through optimised compute.
- Designed 50+ reusable ADF ingestion frameworks, slashing data pipeline development effort and onboarding time by 50%.
- Optimised legacy data pipelines to increase runtime efficiency by 50%.
- Implemented Azure DevOps CI/CD practices that improved data platform release efficiency by 30%.

Project 3: Cloud Data Migration & Automation Platform

Role: Senior Data Engineer

Project Description:

Developed and executed enterprise-scale cloud migration strategies to move on-premises data assets and workloads to Google Cloud Platform (GCP). The project focused on automation, orchestration, and operational optimization to improve scalability, reliability, and maintainability of data processing environments.

Responsibilities:

- Created and implemented automation scripts for efficient migration of on-premises data to Google Cloud (GCP), optimising cloud integration.
- Implemented Autosys orchestration to enhance operational efficiency and minimise risks.

Project 4: Metadata-Driven Spark Processing Framework

Role: Senior Data Engineer

Project Description:

Developed a reusable metadata-driven Spark framework designed to accelerate onboarding of multiple enterprise data sources while standardizing data transformation and enrichment processes. The framework significantly reduced development effort and improved scalability across data processing workflows.

Responsibilities:

- Developed a Generic CCF Enrichment & GLAAS Service framework for metadata-driven Spark pipelines.
- Designed the framework to process and transform raw data according to complex business rules.

- Reduced redundant development efforts for onboarding multiple data sources.
- Improved the overall onboarding process by 35% through framework standardization and automation.

Project 5: Insurance Data Platform Modernization

Role: Data Engineer

Project Description:

Modernized legacy insurance data processing systems by migrating large-scale Mainframe data into a modern Data Lake architecture. The solution improved scalability, automation, validation processes, and data accessibility for downstream analytics and reporting.

Responsibilities:

- Orchestrated 30+ Spark data pipelines for insurance data migration from Mainframe to Data Lake.
- Modernised legacy ELT/ETL processes.
- Developed metadata-driven Mainframe data migration frameworks.
- Automated validation of 2,000+ source files during each extraction cycle.
- Improved migration accuracy, scalability, and operational efficiency through automation.

Project 6: Shipping Management & Carrier Integration Platform

Role: Software Engineer

Project Description:

Developed and enhanced enterprise shipping management solutions by integrating carrier services, freight auditing capabilities, and routing guide functionality. The project focused on improving shipping operations, automation, and long-term software maintainability.

Responsibilities:

- Developed and maintained scalable software solutions in line with business requirements.
- Developed Carrier Integration into the shipping management product TRAX.
- Designed and developed Zebra Printer Label Freight Audit and Routing Guide functionality.
- Developed scalable and maintainable code, ensuring long-term stability of the software.

Project 7: Enterprise Application Development & Support

Role: Software Engineer

Project Description:

Worked on enterprise application development, maintenance, and support activities in a fast-paced environment, contributing to software delivery, issue resolution, and project execution within tight timelines.

Responsibilities:

- Proved successful working within tight deadlines and a fast-paced environment.

STRATEGIC DATA & AI INITIATIVES

Enterprise Lakehouse Implementation

- Engineered a unified Delta Lake architecture optimising data storage and reducing analytical query latency by 35% for cross-functional business users.

Automated Data Quality via Agentic AI

- Architected intelligent, agent-driven workflows embedded directly into enterprise pipelines to automate continuous data quality assurance and ensure high-fidelity data downstream.

Conversational BI Architecture

- Architected an enterprise-grade Conversational BI platform leveraging native Databricks GenAI services, enabling secure, natural language data exploration directly on top of the Lakehouse.